

Keshika Arunkumar

Boston, MA | +1(857) 423 5586 | keshikaarunkumar01@gmail.com | LinkedIn | Github

SUMMARY

Data Engineer with **over 3 years** of experience building **ETL pipelines**, distributed data processing workflows, and analytics platforms across mobility, marketing technology, and enterprise infrastructure domains using **PySpark**, Apache Kafka, Databricks, Snowflake, **AWS Glue**, Hadoop, **SQL**, **dbt**, Airflow, and BigQuery, supporting more than **250 recurring reporting** and ingestion workflows across enterprise data operations while reducing batch processing interruptions and reporting delays.

PROFESSIONAL EXPERIENCE

Data Engineer
Uber

USA

Feb 2026 - Present

- Built data pipelines processing nearly **4M+ ride** and trip events daily using **PySpark** and **Apache Kafka** across **Snowflake** environments, reducing operational reporting delays by 4 hours for marketplace analytics teams.
- Automated **Delta Lake** ingestion workflows using **Databricks**, **Airflow**, and **SQL** validation checks for rider and driver datasets, removing more than **850** recurring weekly data quality exceptions from reporting systems.
- Refined **PostgreSQL** and **BigQuery** query performance through indexing, partitioning, and **ETL optimization** techniques, accelerating **140+** recurring analytics workloads supporting operations and financial reporting teams.
- Standardized **Terraform** automation and **CI/CD deployment** processes for containerized analytics services, reducing manual release validation effort by **9 hours** across monthly production deployments.
- Coordinated **data validation initiatives** with product managers, analysts, and backend engineers using **Python** and **dbt** monitoring models, resolving **320+ schema** and transformation issues before executive dashboard releases.

Data Engineer
Epsilon

USA

Jan 2025 - Aug 2025

- Engineered customer segmentation pipelines using **Spark SQL** and **Amazon Redshift** for multi-channel campaign data, processing nearly **3.8M** customer engagement records during quarterly targeting cycles.
- Consolidated batch ingestion processes through **AWS Glue**, **Python**, and **Medallion Architecture** transformation layers, reducing recurring advertising data load failures by nearly 120 records across weekly audience delivery runs.
- Integrated event-stream pipelines with **Apache NiFi** and **MongoDB** for loyalty and engagement platforms, reducing turnaround time for campaign attribution reporting across **6 regional** marketing programs.
- Migrated legacy ETL mappings into **Informatica PowerCenter** and **Oracle SQL** pipelines, cutting nightly customer analytics refresh windows by nearly **3 operational hours** across production reporting environments.
- Validated audience and conversion reporting layers using **dbt**, **Tableau**, and automated reconciliation scripts, identifying more than **90** mapping inconsistencies before downstream media performance reviews.

Associate Data Engineer
Dell Technologies

IND

Jul 2021 - Aug 2023

- Developed server log ingestion pipelines using **Hadoop**, **Hive**, and **Sqoop** for infrastructure monitoring data, processing nearly **1.6M system** records supporting weekly storage and server performance reviews.
- Administered ETL scheduling routines through **Unix Shell Scripting** and **Control-M**, reducing delayed batch execution incidents across 45 recurring monthly processing jobs within production environments.
- Transformed hardware inventory feeds using **PySpark** and **HDFS** staging layers, improving reconciliation coverage across **12 enterprise** asset tracking reports used by infrastructure support teams.
- Optimized transactional queries within **MySQL** and **SQL Server** reporting environments, shortening dashboard refresh intervals by nearly **90 minutes** for supply chain analytics and inventory operations users.
- Streamlined data extraction procedures through **Talend** and **XML** parsing workflows, reducing manual remediation effort tied to more than 80 malformed vendor records during quarterly compliance audits.
- Monitored production data validation checks using **Power BI**, **Excel**, and automated exception tracking scripts, preventing nearly **60** reporting discrepancies from reaching infrastructure compliance review cycles.

TECHNICAL SKILLS

- Programming Languages:** Python (Pandas, NumPy, PySpark), SQL, T-SQL, R, Unix Shell Scripting
- Big Data & Data Engineering:** Apache Kafka, Hadoop, Hive, Sqoop, HDFS, ETL/ELT Pipelines, Delta Lake, Batch Processing, Event-Stream Processing, Data Validation, Data Reconciliation, Data Quality, Query Optimization, Ingestion Frameworks
- Data Modeling & Warehousing:** Star Schema, Dimensional Modeling, SCD Type 1, SCD Type 2, Medallion Architecture
- Cloud & Lakehouse Platforms:** AWS (S3, EKS), Azure (Data Factory, SQL Database, ADLS Gen2), GCP, Microsoft Fabric (OneLake, Lakehouse), Databricks, Snowflake, Amazon Redshift, BigQuery
- Databases & Warehouses:** PostgreSQL, MySQL, SQL Server, Oracle SQL, MongoDB

- **ETL & Orchestration Tools:** Apache Airflow, Azure Data Factory, Informatica PowerCenter, Talend, dbt, Great Expectations
 - **CI/CD & DevOps Tools:** Terraform, Docker, Kubernetes, Git, GitHub Actions, CI/CD Pipelines
 - **Visualization & Reporting Tools:** Power BI, Tableau, Excel
 - **Data Formats & Processing:** XML, Indexing, Partitioning
-

EDUCATION

Northeastern University | Master of Science in Data Analytics Engineering | Boston, MA |

Dec 2025

PROJECTS

RideStream: Real-Time Uber Data Pipeline

- Constructed real-time ingestion pipelines using **Azure Event Hubs**, **PySpark Structured Streaming**, and **Databricks**, processing more than **10K ride** events into Bronze, Silver, and Gold Delta Lake layers.
- Modeled analytical datasets through **Star Schema**, **SCD Type 1 & 2**, and Medallion Architecture design patterns, reducing query joins across downstream reporting workflows.

OutbreakLens: Global Disease Surveillance Pipeline

- Deployed automated ingestion workflows through **Apache Airflow** and **PostgreSQL**, loading more than 260K public health records from 3 external data sources through 5 scheduled DAGs.
- Curated transformation models using **dbt** and automated data quality validations, executing 27 validation checks supporting outbreak trend monitoring and reporting accuracy.

Financial Data Quality Monitor | SQL | Airflow | Great Expectations

- Built automated data quality pipelines using Airflow, SQL, and Great Expectations, validating 5.6K+ financial records across multiple data sources.
 - Developed anomaly detection and monitoring workflows with 29 automated validation checks, identifying 230+ data quality issues through reconciliation and health scoring.
-

CERTIFICATIONS

- **Microsoft Certified: Fabric Data Engineer Associate**
- **Databricks – Generative AI Application Development**
- **Databricks – GenAI Application Evaluation and Governance**
- **Databricks – Generative AI Application Deployment and Monitoring**
- **Google Data Analytics**
- **Tableau Certified (Simplilearn)**
- **AWS Introduction to Generative AI**